

▼ Assignment 4A - Vector Addition

```

!nvcc --version

nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2022 NVIDIA Corporation
Built on Wed_Sep_21_10:33:58_PDT_2022
Cuda compilation tools, release 11.8, V11.8.89
Build cuda_11.8.r11.8/compiler.31833905_0

code = """

#include <stdio.h>
#include <stdlib.h>

#define N 1000000

__global__ void add(int *a, int *b, int *c) {
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    if (tid < N) {
        c[tid] = a[tid] + b[tid];
    }
}

int main() {
    int *a, *b, *c;
    int *d_a, *d_b, *d_c;
    int size = N * sizeof(int);

    // Allocate memory on host
    a = (int*)malloc(size);
    b = (int*)malloc(size);
    c = (int*)malloc(size);

    // Initialize arrays
    for (int i = 0; i < N; i++) {
        a[i] = i;
        b[i] = i * 2;
    }

    // Allocate memory on device
    cudaMalloc(&d_a, size);
    cudaMalloc(&d_b, size);
    cudaMalloc(&d_c, size);

    // Copy data from host to device
    cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);

    // Launch kernel with 1 million threads
    add<<<(N + 255) / 256, 256>>>(&d_a, &d_b, &d_c);

    // Copy result from device to host
    cudaMemcpy(c, &d_c, size, cudaMemcpyDeviceToHost);

    // Print first and last elements of result
    printf("c[0]=%d, c[%d] = %d", c[0], N-1, c[N-1]);

    // Free memory
    free(a);
    free(b);
    free(c);
    cudaFree(&d_a);
    cudaFree(&d_b);
    cudaFree(&d_c);

    return 0;
}
"""

text_file = open("assign4.cu", "w")
text_file.write(code)
text_file.close()

!nvcc assign4.cu

!./a.out

c[0]=0, c[999999] = 2999997

!nvprof ./a.out

```

==1065== NVPROF is profiling process 1065, command: ./a.out
 ==1065== Profiling application: ./a.out
 ==1065== Profiling result:

	Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:		51.91%	1.9264ms	1	1.9264ms	1.9264ms	1.9264ms	[CUDA memcpy DtoH]
		46.81%	1.7372ms	2	868.62us	813.34us	923.90us	[CUDA memcpy HtoD]
		1.28%	47.455us	1	47.455us	47.455us	47.455us	add(int*, int*, int*)
API calls:		97.16%	268.79ms	3	89.598ms	100.42us	268.58ms	cudaMalloc
		2.08%	5.7679ms	3	1.9226ms	1.0787ms	3.5174ms	cudaMemcpy
		0.44%	1.2053ms	1	1.2053ms	1.2053ms	1.2053ms	cuDeviceGetPCIBusId
		0.24%	676.53us	3	225.51us	205.02us	237.71us	cudaFree
		0.05%	143.16us	101	1.4170us	133ns	73.062us	cuDeviceGetAttribute
		0.01%	37.564us	1	37.564us	37.564us	37.564us	cudaLaunchKernel
		0.01%	25.957us	1	25.957us	25.957us	25.957us	cuDeviceGetName
		0.00%	1.6580us	3	552ns	218ns	1.2050us	cuDeviceGetCount
		0.00%	1.0080us	2	504ns	192ns	816ns	cuDeviceGet
		0.00%	505ns	1	505ns	505ns	505ns	cuModuleGetLoadingMode
		0.00%	447ns	1	447ns	447ns	447ns	cuDeviceTotalMem
		0.00%	280ns	1	280ns	280ns	280ns	cuDeviceGetUuid
c[0]=0, c[999999] = 2999997								