

```

!nvcc --version

nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2022 NVIDIA Corporation
Built on Wed_Sep_21_10:33:58_PDT_2022
Cuda compilation tools, release 11.8, V11.8.89
Build cuda_11.8.r11.8/compiler.31833905_0

code = """
#include <stdio.h>
#include <stdlib.h>

#define N 1024
#define BLOCK_SIZE 16

__global__ void matrixMul(int *a, int *b, int *c, int width) {
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;
    int sum = 0;
    for (int i = 0; i < width; i++) {
        sum += a[row * width + i] * b[i * width + col];
    }
    c[row * width + col] = sum;
}

int main() {
    int *a, *b, *c;
    int *d_a, *d_b, *d_c;
    int size = N * N * sizeof(int);

    // Allocate memory on host
    a = (int*)malloc(size);
    b = (int*)malloc(size);
    c = (int*)malloc(size);

    // Initialize matrices
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            a[i * N + j] = i + j;
            b[i * N + j] = i - j;
        }
    }

    // Allocate memory on device
    cudaMalloc(&d_a, size);
    cudaMalloc(&d_b, size);
    cudaMalloc(&d_c, size);

    // Copy data from host to device
    cudaMemcpy(d_a, a, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_b, b, size, cudaMemcpyHostToDevice);

    // Launch kernel with 2D grid and 2D block
    dim3 dimBlock(BLOCK_SIZE, BLOCK_SIZE);
    dim3 dimGrid((N + dimBlock.x - 1) / dimBlock.x, (N + dimBlock.y - 1) / dimBlock.y);
    matrixMul<<<dimGrid, dimBlock>>>(d_a, d_b, d_c, N);

    // Copy result from device to host
    cudaMemcpy(c, d_c, size, cudaMemcpyDeviceToHost);

    // Print first and last elements of result
    printf("c[0][0] = %d, c[%d][%d] = %d", c[0], N-1, N-1, c[(N-1) * N + (N-1)]);

    // Free memory
    free(a);
    free(b);
    free(c);
    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(d_c);

    return 0;
}

"""

text_file = open("assign4b.cu", "w")
text_file.write(code)
text_file.close()

!nvcc assign4b.cu

!./a.out

c[0][0] = 357389824, c[1023][1023] = -714255872

```

```
!nvcc -o a.out
```

```
nvprof ./a.out
==939== NVPROF is profiling process 939, command: ./a.out
==939== Profiling application: ./a.out
c[0][0] = 357389824, c[1023][1023] = -714255872==939== Profiling result:
      Type  Time(%)    Time       Calls       Avg       Min       Max  Name
GPU activities:  70.94%   9.1100ms         1   9.1100ms   9.1100ms   9.1100ms  matrixMul(int*, int*, int*, int)
                15.21%   1.9532ms         1   1.9532ms   1.9532ms   1.9532ms  [CUDA memcpy DtoH]
                13.85%   1.7782ms         2   889.09us   884.09us   894.08us  [CUDA memcpy HtoD]
API calls:      94.19%  276.44ms         3   92.146ms  119.31us  276.20ms  cudaMalloc
                5.11%   15.010ms         3   5.0034ms   1.1188ms  12.758ms  cudaMemcpy
                0.35%   1.0392ms         1   1.0392ms   1.0392ms   1.0392ms  cuDeviceGetPCIBusId
                0.23%   678.78us         3   226.26us   219.79us   233.19us  cudaFree
                0.07%   212.76us        101   2.1060us         310ns   78.656us  cuDeviceGetAttribute
                0.02%   56.497us         1   56.497us   56.497us   56.497us  cuDeviceGetName
                0.01%   40.948us         1   40.948us   40.948us   40.948us  cudaLaunchKernel
                0.00%   2.2660us         3     755ns         377ns   1.5050us  cuDeviceGetCount
                0.00%   1.2130us         2     606ns         306ns         907ns  cuDeviceGet
                0.00%      857ns         1     857ns         857ns         857ns  cuDeviceTotalMem
                0.00%      466ns         1      466ns         466ns         466ns  cuModuleGetLoadingMode
                0.00%      440ns         1      440ns         440ns         440ns  cuDeviceGetUuid
```

[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 3:56 PM

